Deploying LLaMA-2-7B on a Raspberry Pi 4

A Four-Stage Compression Pipeline for On-Device LLM Inference

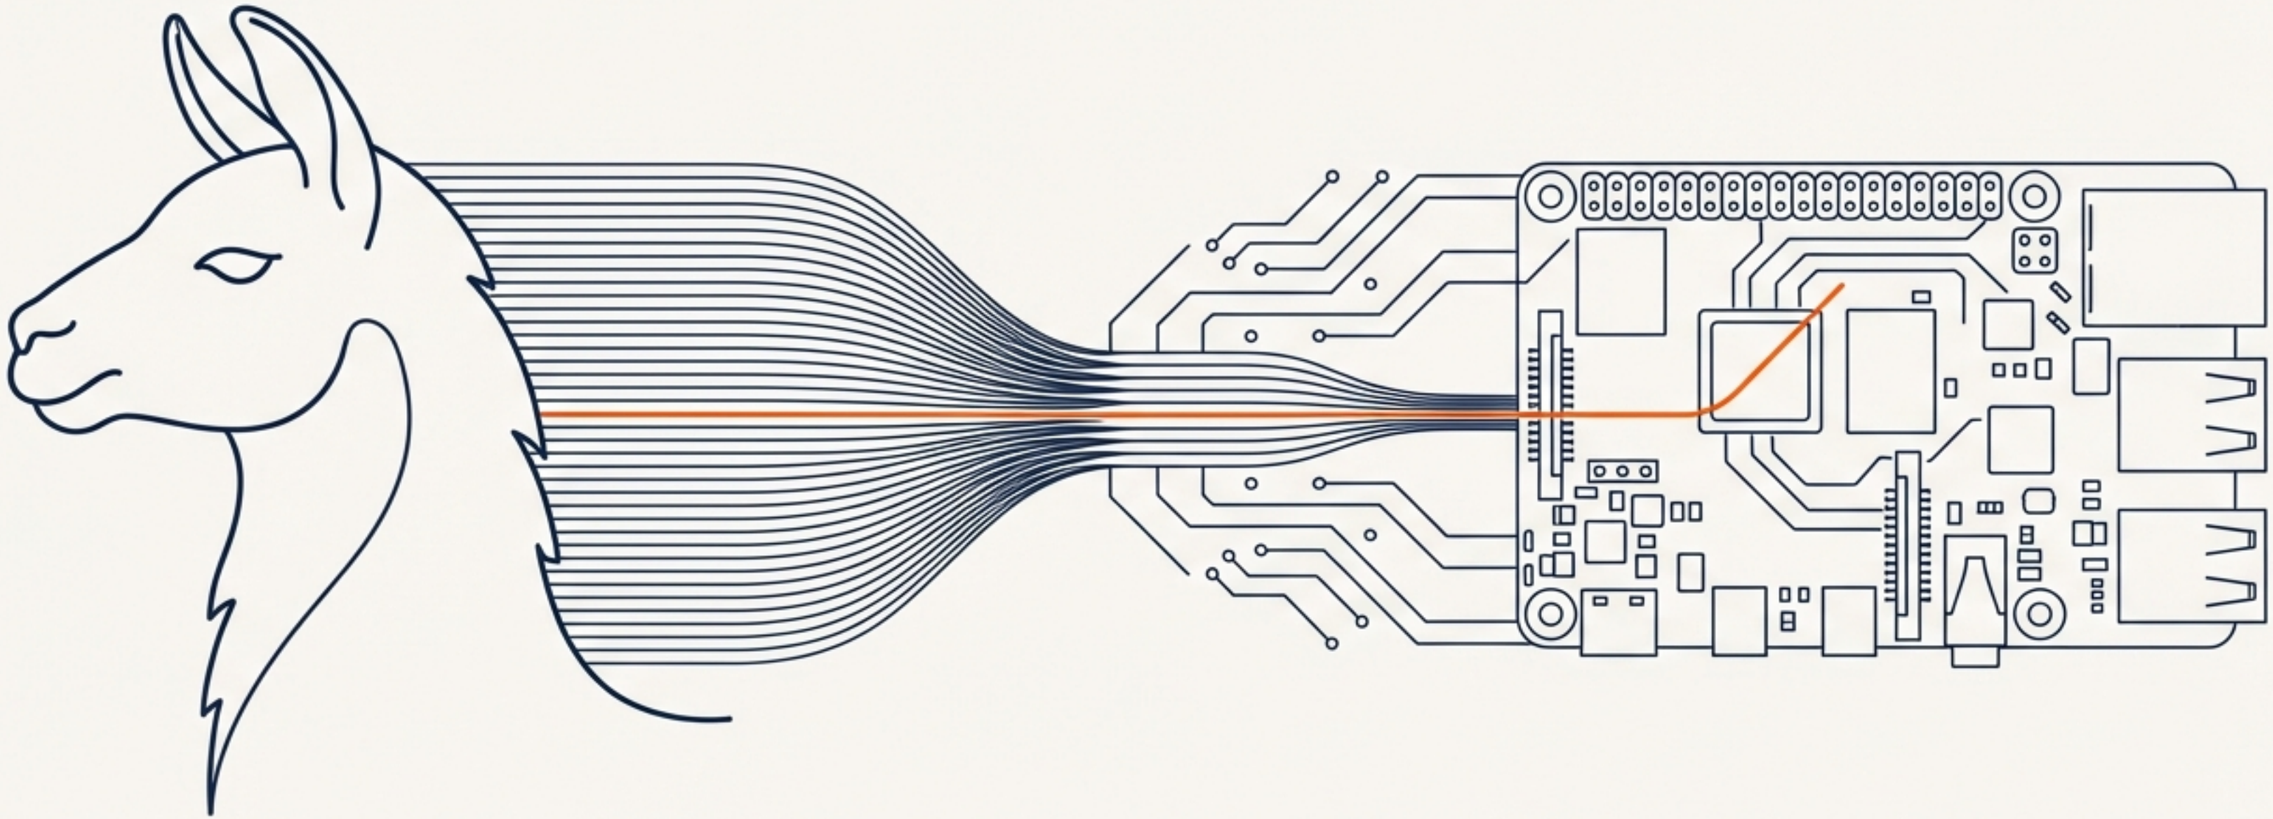

MODEL	SIZE REDUCTION	TARGET DEVICE
LLaMA-2-7B	13.5 GB → 3.29 GB (75.6%)	Raspberry Pi 4 (4GB RAM)

The New Frontier: Moving Large Language Models to the Edge

The Cloud Confinement

Large Language Models like LLaMA-2 have transformed AI, but their immense size (e.g., 13.5 GB for LLaMA-2-7B in FP16) restricts them to powerful cloud servers.

The Edge Imperative

Deploying models on-device is critical for:

- **Privacy**: User data never leaves the local environment.
- **Latency**: Eliminates network round-trip delays for critical applications.
- **Accessibility**: Enables offiine use and democratizes AI without reliance on cloud infrastructure.

The Challenge

The target environment—a Raspberry Pi 4 with 4GB RAM—is a proxy for countless resource-constrained devices. It represents a hostile environment for multi-billion parameter models.

VISUAL DIAGRAM

Cloud-Based

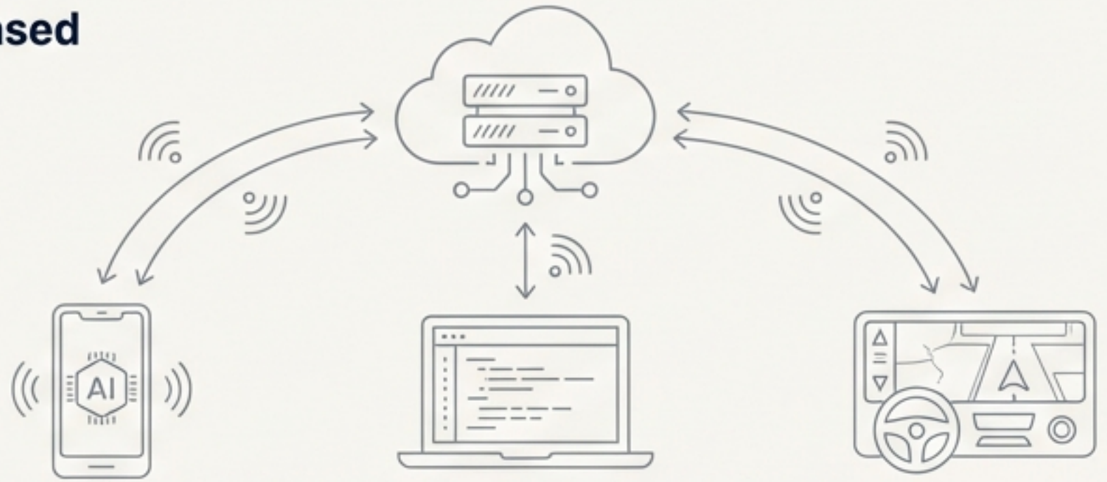

Network Dependency: Data Transfer & Latency

On-Device

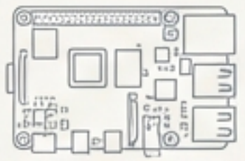

Representative Target Hardware

Raspberry Pi 4

Self-Contained Processing: Local Inference, No Network Required

The Mission: Fit a 13.5 GB Model into a 4 GB Device

The central challenge is to compress the 6.74 billion parameter LLaMA-2 model to operate within the 4 GB RAM limit of a Raspberry Pi 4, while preserving coherent text generation.

Original Model (13.5 GB)

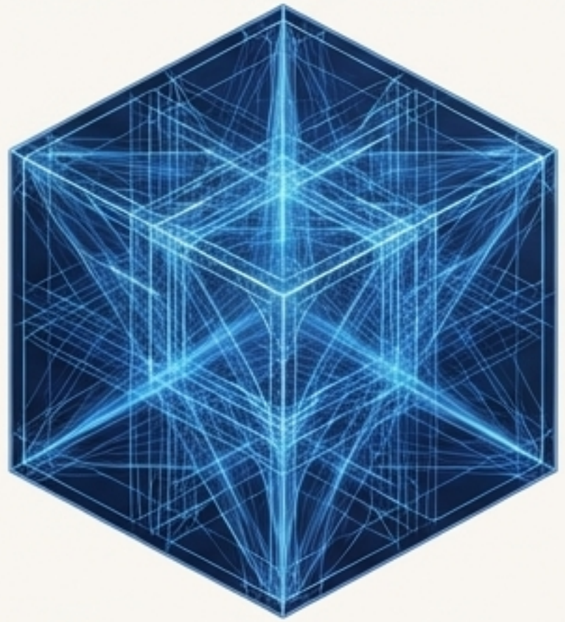

Compression Pipeline

Target Memory Footprint (< 4 GB)

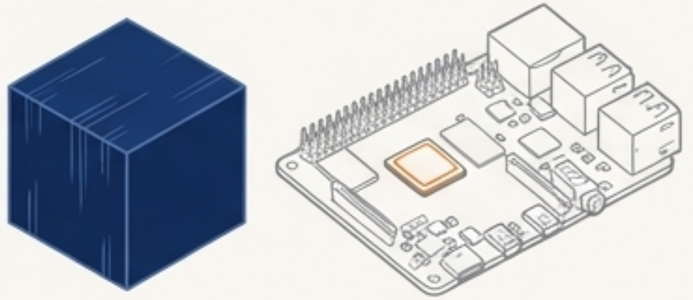

Key Constraints

Initial Model Size:	13.5 GB (FP16)
Total Parameters:	6.74 Billion
Target Device RAM:	4 GB
Core Requirement:	Must maintain acceptable text generation quality without fine-tuning.

The Strategy: A Four-Stage Compression Pipeline

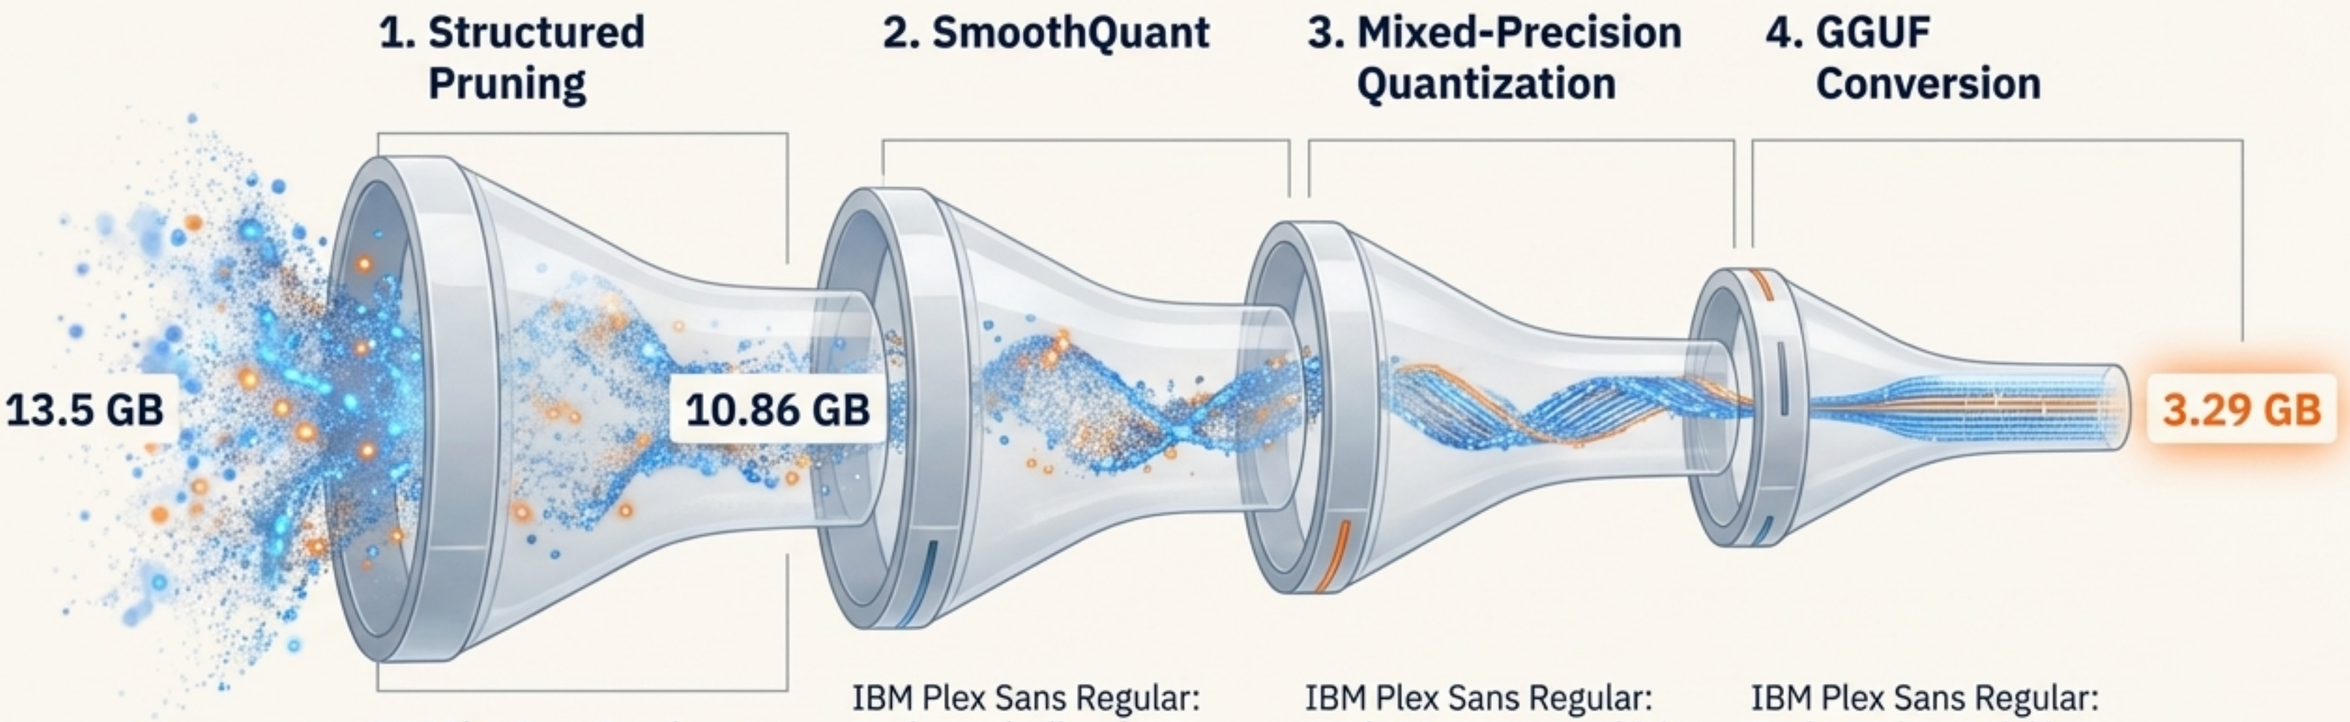

IBM Plex Sans Regular: Remove entire neurons to reduce model complexity while maintaining a dense structure for efficient CPU computation.

IBM Plex Sans Regular: Mathematically migrate quantization difficulty from activations to weights, enabling more aggressive compression without quality collapse.

IBM Plex Sans Regular: Apply a custom sensitivity analysis to allocate higher higher bit-widths to critical layers and lower bit-widths to robust ones.

IBM Plex Sans Regular: Package the model into an optimized format for efficient, memory-mapped loading and CPU inference.

Phase 1: Reducing Model Complexity with Structured Pruning

Why Structured Pruning?

Unstructured pruning creates sparse matrices requiring specialized hardware. For edge CPUs, structured pruning is the only viable option as it produces smaller, dense matrices compatible with standard hardware.

Methodology

- **Target:** Feed-Forward Networks (FFNs), which constitute ~67% of the model's parameters.
- **Technique:** Activation-based importance scoring. Neurons are ranked by their mean absolute activation on calibration data (128 samples from WikiText-2).
- **Action:** Removed the bottom 20% of neurons, reducing the FFN intermediate dimension from 11,008 to 8,704.

Results of Phase 1

Metric	Before Pruning	After Pruning
Total Parameters	6.74B	5.83B (-13.4%)
Model Size (FP16)	13.50 GB	10.86 GB
Perplexity (WikiText-2)	8.47	11.40 (+34.5%)

Key Takeaway: A 13.4% parameter reduction was achieved with an expected perplexity increase. Qualitative evaluation confirmed the model remained coherent.

The Critical Insight: Correctly Implementing SmoothQuant

The Problem: Activation Outliers

Naive quantization fails because transformer activations have outliers—channels with values 10-100x larger than the median. This forces a choice between clipping critical information or losing precision for typical values.

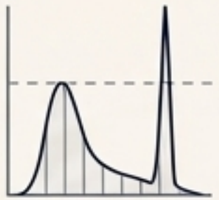

The Solution: SmoothQuant

SmoothQuant migrates quantization difficulty from activations to weights via a per-channel scaling factor `s. The key is preserving mathematical equivalence.

The Implementation Detail That Matters

To maintain equivalence, scaling must be applied to both the preceding LayerNorm and the subsequent Linear layer. This is a commonly overlooked detail.

Correct Implementation (Mathematical Equivalence)

$Y = (X / s) * (s * W) = X * W$

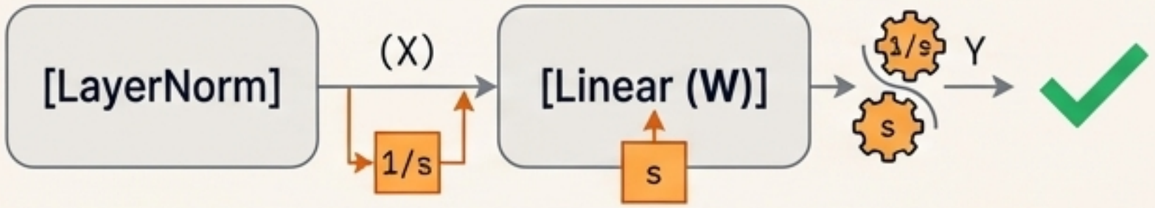

Incorrect Implementation (Scaled Output)

$Y_{wrong} = X * (s * W) \neq X * W$

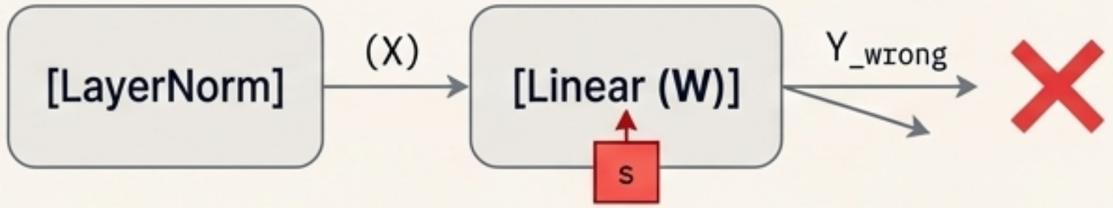

Our initial implementation made this mistake, producing completely incoherent outputs. Rigorous mathematical verification is essential.

Phase 3: Optimizing Bit-Width with Custom Sensitivity Analysis

Motivation

Not all layers are equally sensitive to quantization. Mixed-precision allocates bits intelligently, assigning higher precision to sensitive layers and lower precision to robust ones.

Custom Sensitivity Framework

We developed a lightweight measure combining weight and activation statistics to rank layer sensitivity.

Weight-Based Score

$$S_{weight} = \sigma_W * (1 + 10 * r_{outlier})$$

Measures standard deviation and the fraction of weights >3 std devs from the mean.

Activation-Based Score

$$S_{activation} = \sigma_A * (1 + \max(|A|)/100)$$

Measures activation standard deviation and maximum outlier value from calibration.

Combined Score

$$S_{combined} = 0.7 * S_{weight} + 0.3 * S_{activation}$$

Bit-Width Allocation Strategy

Layers were ranked by their combined score, leading to this allocation:

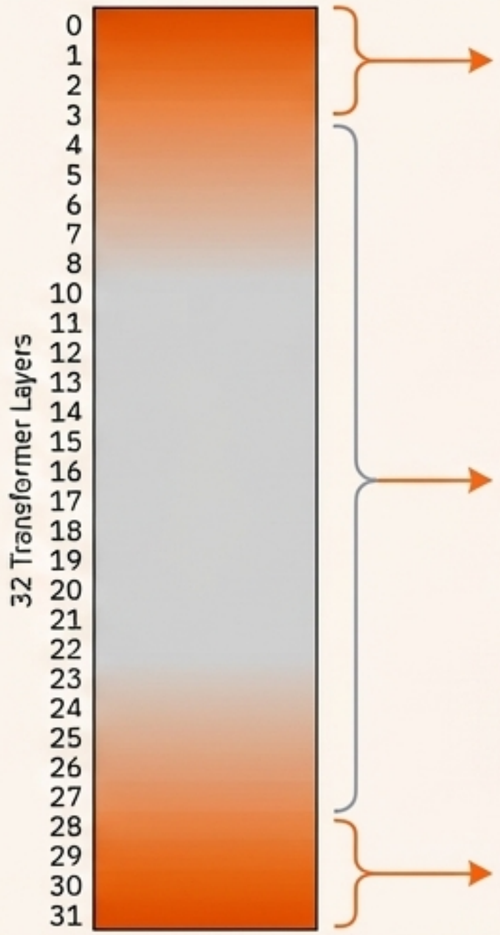

Assigned 8-bit quantization (Q8)

Primarily early (0-3) and late (28-31) transformer layers, especially attention value and FFN down projections.

Assigned 4-bit quantization (Q4)

Middle layers (10-20) and FFN gate/up projections proved most robust.

Mission Success: Results Dashboard

Final Compression

75.6%

% COMPRESSION

Model Configuration	Size (GB)	Fits 4GB RAM?
Original LLaMA-2-7B (FP16)	13.50	**No**
Pruned + Naive Q4_0	3.09	Yes
Pruned + SmoothQuant + Q4_K_M	**3.29**	**Yes**

Final model is 3.29 GB, leaving ~700MB for OS and KV Cache on the Pi.

Inference Performance (CPU)

10.5

tok/s (Decode)

Metric	Value
Time to First Token	279 ms
Prefill Throughput	21.5 tok/s
Decode Throughput	10.5 tok/s

Projected throughput on Raspberry PI 4 hardware: 2-3 tokens/second.

Language Quality

Coherent Text Generation Maintained

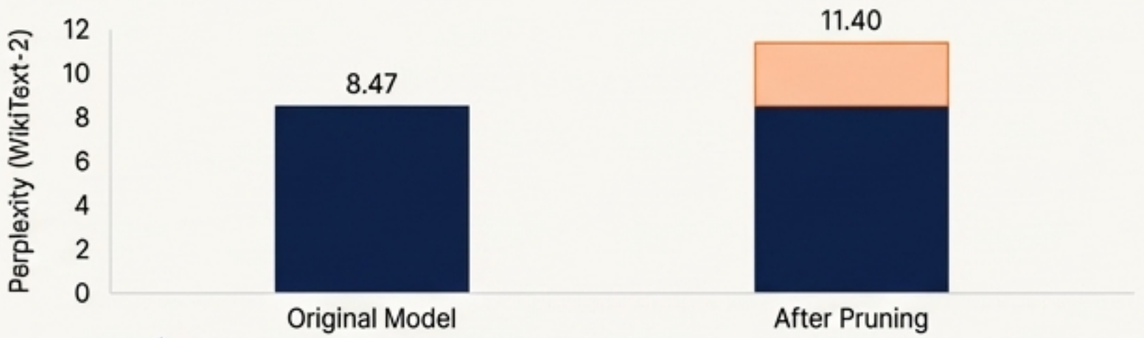

The primary quality degradation came from pruning. SmoothQuant and quantization had minimal additional impact.

Evidence of Coherence: The Compressed Model in Action

Quantitative metrics confirm success, but the ultimate test is qualitative. The final Q4_K_M model demonstrates it retains core language modeling capabilities.

PROMPT:
`` `The capital of Japan is` ``

MODEL OUTPUT:
"The capital of Japan is one of the most expensive cities in the world. Its GDP per capita is over $86,000, which makes Tokyo the 2nd richest city in the world."

ANALYSIS:
Coherent, factually plausible, and grammatically correct.

PROMPT:
`` `Artificial intelligence is` ``

MODEL OUTPUT:
"Artificial intelligence is going to become a part of our daily lives in the next decade. This is the the belief of the co-founder of Google DeepMind."

ANALYSIS:
Demonstrates logical flow, contextual relevance, and proper sentence structure.

The outputs confirm the compressed model is not just small, but also practically useful for text generation.

The Debrief: Key Insights and Lessons Learned

Technique	Primary Contribution
Structured Pruning	**13.4% parameter reduction** without fine-tuning.
SmoothQuant	Critical for quality preservation, enabling aggressive 4-bit quantization.
Sensitivity Analysis	Provided theoretical justification for mixed-precision strategy.
Q4_K_M Quantization	Final **75.6% size reduction** to meet the memory target.

The SmoothQuant Pitfall

The absolute necessity of scaling both LayerNorm and Linear weights to maintain mathematical equivalence was the most critical learning. A seemingly correct implementation that only scales linear weights will fail **catastrophically**. This highlights the importance of deep architectural understanding in quantization.

A sequential pipeline where each stage prepares the model for the next is highly effective for post-training compression.

Debrief Continued: Limitations and Future Missions

Acknowledged Limitations

- **No Fine-Tuning:** The 34.5% perplexity increase from pruning could likely be recovered via knowledge distillation or further training, which was beyond the project's scope.
- **Inference Speed:** At a projected 2-3 tokens/sec on Raspberry Pi, the model is suitable for batch tasks, but not interactive chat.
- **Generalization:** Evaluation was limited to perplexity and qualitative generation; performance on specific downstream tasks (e.g., Q&A) was not measured.

Future Missions (Future Work)

1. **Recover Quality:** Use knowledge distillation with the original model as a teacher to regain perplexity lost during pruning.
2. **Accelerate Inference:** Implement speculative decoding with a smaller draft model to improve token generation speed.
3. **Further Compression:** Explore attention head pruning as an orthogonal technique to FFN pruning.
4. **Hardware Optimization:** Leverage the Raspberry Pi's VideoCore GPU for hardware-specific acceleration.

Conclusion: Billion-Parameter Language Models are Ready for the Edge

Summary of Achievement

This work successfully demonstrated a practical pipeline to compress LLaMA-2-7B by **75.6%**, reducing its size from 13.5 GB to 3.29 GB. The compressed model operates effectively within the 4GB RAM constraint of a Raspberry Pi 4 while maintaining coherent language capabilities.

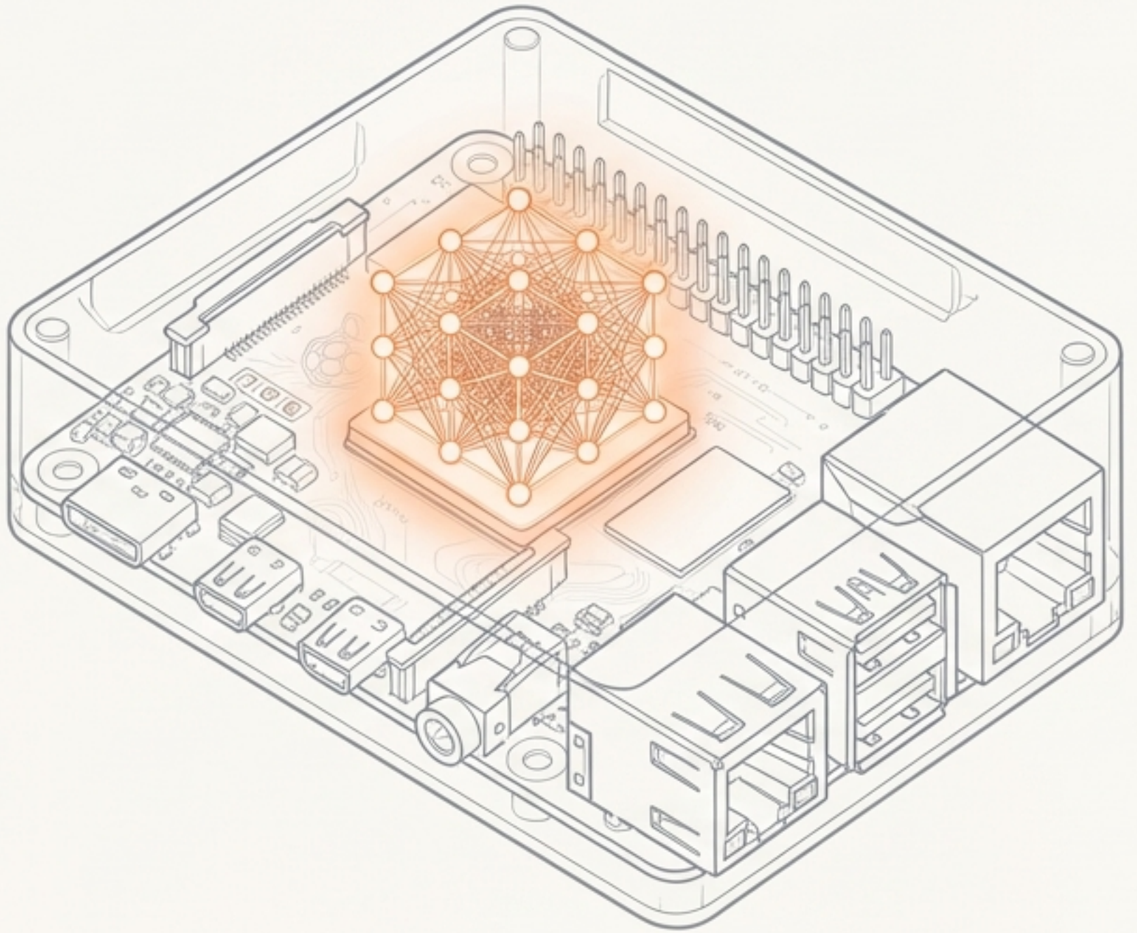

Key Contributions Recap

- A robust four-stage compression methodology.
- A validated, correct implementation of SmoothQuant.
- A custom sensitivity framework justifying mixed-precision quantization.

The Broader Impact

The success of this mission proves that with the right combination of techniques, it is possible to deploy powerful foundation models on widely available, consumer-grade hardware without specialized accelerators. This unlocks a new class of applications that are privacy-preserving, low-latency, and perpetually available, regardless of network connectivity.